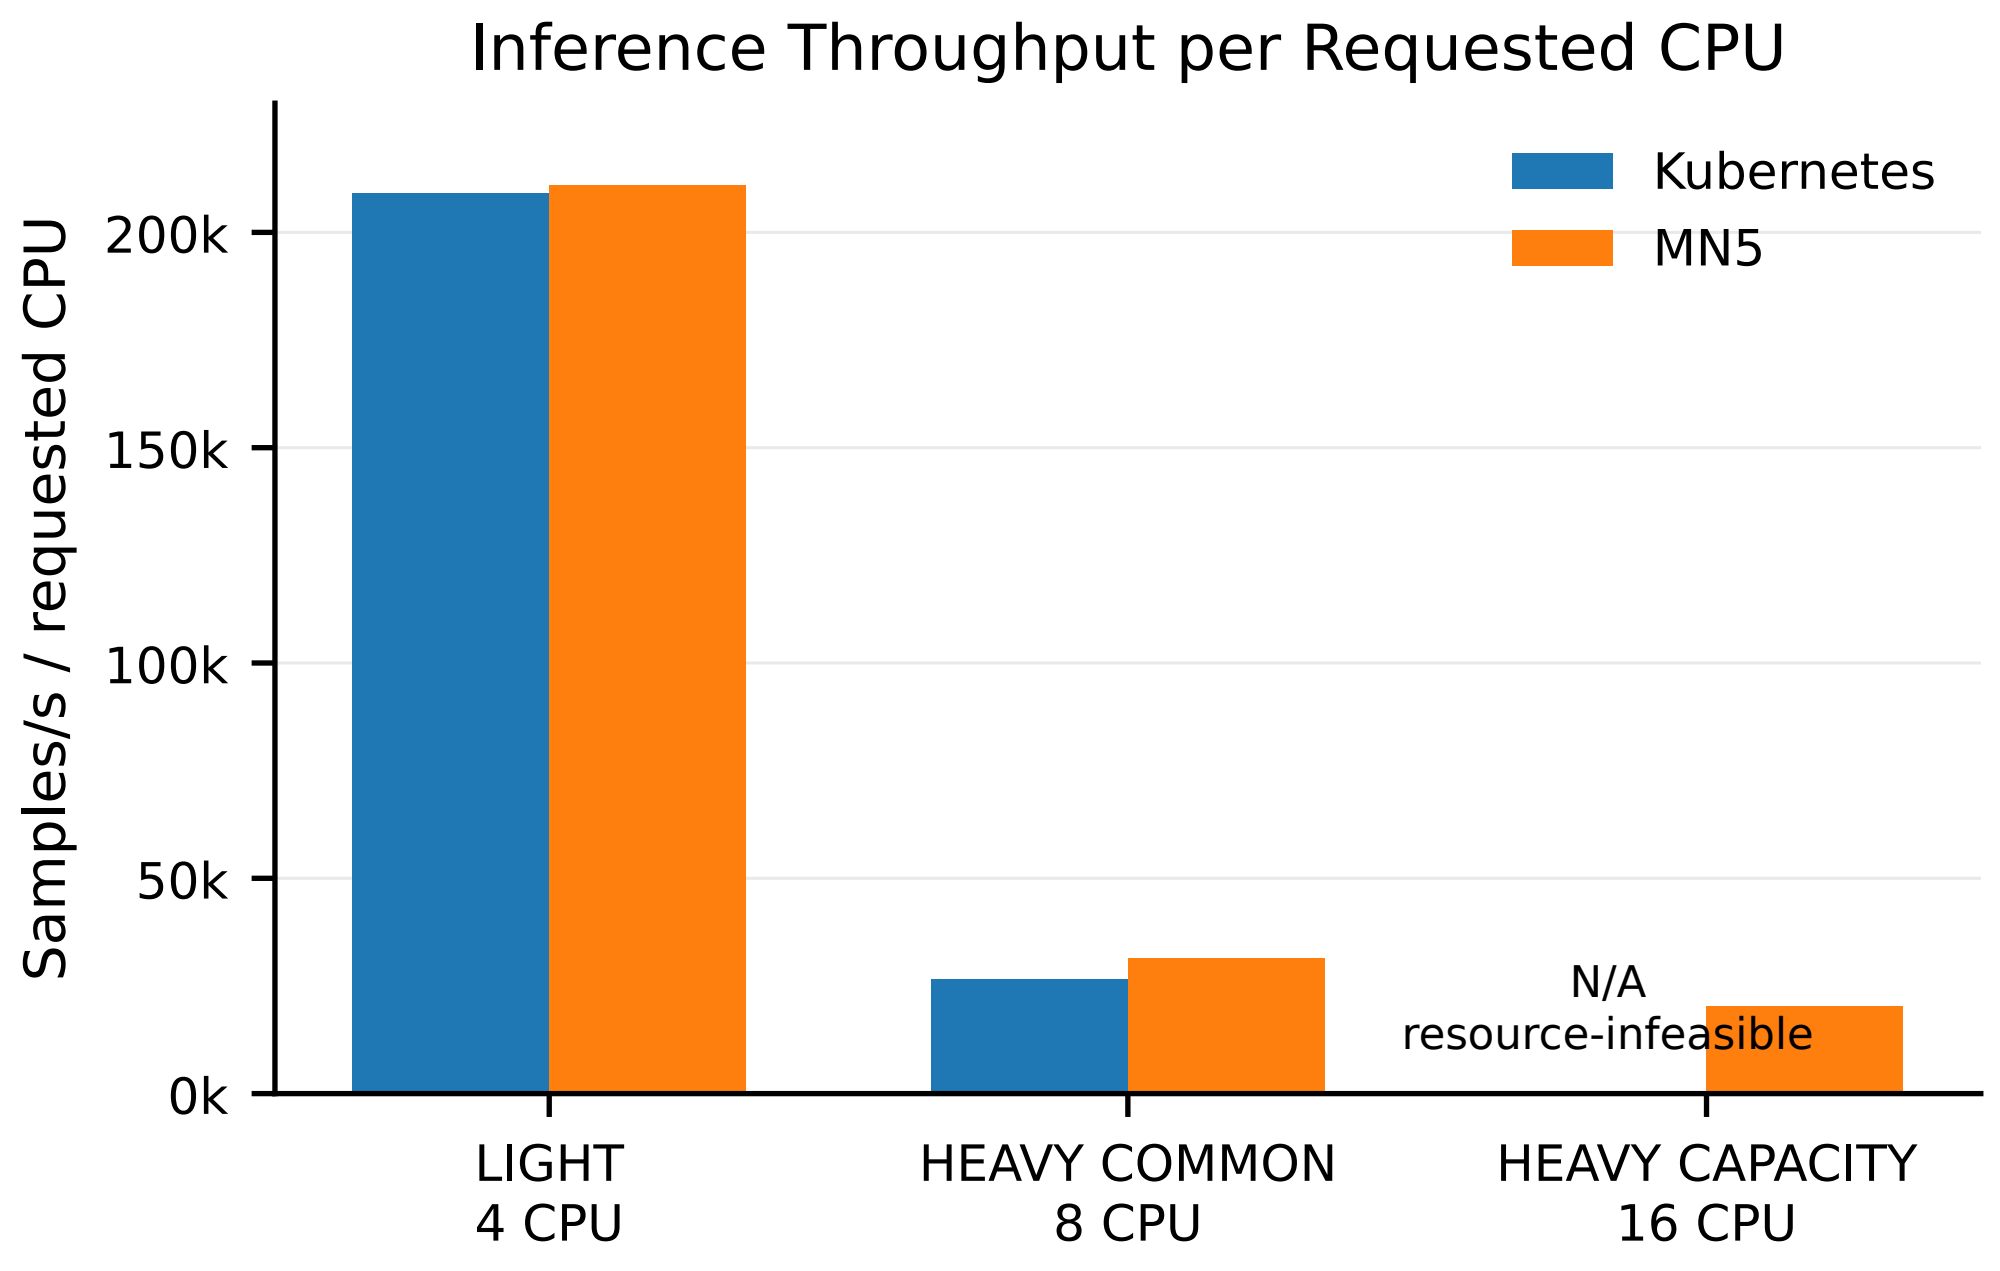
Inference Throughput per Requested CPU
Kubernetes
MN5
Samples/s / requested CPU
200k
150k
100k
50k
0k
N/A
resource-infeasible
LIGHT
4 CPU
HEAVY COMMON
8 CPU
HEAVY CAPACITY
16 CPU